



Latency Breakdown at Concurrency=32 (averaged across workloads)

TP-only (2 GPU) EP-only (2 GPU) TP+EP Hybrid (4 GPU)

Mean TTFT



Mean Inter-Token Latency



Mean End-to-End Latency

